

HW2_Submission

May 6, 2024

0.1 Cart Pole Swingup Constrained

```
[ ]: """
    Solution code for the problem "Cart-pole balance".

    Autonomous Systems Lab (ASL), Stanford University
    """

import numpy as np
from scipy.integrate import odeint
import matplotlib.pyplot as plt

from animations import animate_cartpole

import pdb

# Constants
n = 4 # state dimension
m = 1 # control dimension
mp = 2.0 # pendulum mass
mc = 10.0 # cart mass
L = 1.0 # pendulum length
g = 9.81 # gravitational acceleration
dt = 0.1 # discretization time step
animate = False # whether or not to animate results

def cartpole(s: np.ndarray, u: np.ndarray) -> np.ndarray:
    """Compute the cart-pole state derivative

    Args:
        s (np.ndarray): The cartpole state: [x, theta, x_dot, theta_dot], shape (n,)
        u (np.ndarray): The cartpole control: [F_x], shape (m,)

    Returns:
        np.ndarray: The state derivative, shape (n,)
    """
```

```

x, , dx, d = s
sin, cos = np.sin(), np.cos()
h = mc + mp * (sin**2)
ds = np.array(
    [
        dx,
        d,
        (mp * sin * (L * (d**2) + g * cos) + u[0]) / h,
        -((mc + mp) * g * sin + mp * L * (d**2) * sin * cos + u[0] * cos)
        / (h * L),
    ]
)
return ds

def reference(t: float) -> np.ndarray:
    """Compute the reference state ( $\bar{s}$ ) at time  $t$ 

    Args:
        t (float): Evaluation time

    Returns:
        np.ndarray: Reference state, shape (n,)
    """
    a = 10.0 # Amplitude
    T = 10.0 # Period
    # breakpoint()
    # PART (d) #####
    # INSTRUCTIONS: Compute the reference state for a given time
    # raise NotImplementedError()
    return np.array([a * np.sin(2*np.pi*t/T), np.pi, 2*np.pi*a/T * np.cos(2*np.
    pi*t/T), 0]).T
    # END PART (d) #####

def ricatti_recursion(
    A: np.ndarray, B: np.ndarray, Q: np.ndarray, R: np.ndarray
) -> np.ndarray:
    """Compute the gain matrix  $K$  through Ricatti recursion

    Args:
        A (np.ndarray): Dynamics matrix, shape (n, n)
        B (np.ndarray): Controls matrix, shape (n, m)
        Q (np.ndarray): State cost matrix, shape (n, n)
        R (np.ndarray): Control cost matrix, shape (m, m)

    Returns:

```

```

        np.ndarray: Gain matrix K, shape (m, n)
    """
    eps = 1e-4 # Riccati recursion convergence tolerance
    max_iters = 1000 # Riccati recursion maximum number of iterations
    P_prev = np.zeros((n, n)) # initialization
    converged = False
    for i in range(max_iters):
        # PART (b) #####
        # INSTRUCTIONS: Apply the Riccati equation until convergence
        K = -np.linalg.inv(R + B.T @ P_prev @ B) @ B.T @ P_prev @ A
        P_k = Q + A.T @ P_prev @ (A + B @ K)

        # termination condition
        if np.max(np.abs(P_prev-P_k)) < 1e-4:
            converged = True
            break
        else:
            P_prev = P_k
        # END PART (b) #####
    if not converged:
        raise RuntimeError("Riccati recursion did not converge!")
    print("K:", K)
    return K

def simulate(
    t: np.ndarray, s_ref: np.ndarray, u_ref: np.ndarray, s0: np.ndarray, K: np.
    ↪ ndarray
) -> tuple[np.ndarray, np.ndarray]:
    """Simulate the cartpole

    Args:
        t (np.ndarray): Evaluation times, shape (num_timesteps,)
        s_ref (np.ndarray): Reference state  $\bar{s}$ , evaluated at each time  $t$ .
    ↪ Shape (num_timesteps, n)
        u_ref (np.ndarray): Reference control  $\bar{u}$ , shape (m,)
        s0 (np.ndarray): Initial state, shape (n,)
        K (np.ndarray): Feedback gain matrix (Riccati recursion result), shape
    ↪ (m, n)

    Returns:
        tuple[np.ndarray, np.ndarray]: Tuple of:
            np.ndarray: The state history, shape (num_timesteps, n)
            np.ndarray: The control history, shape (num_timesteps, m)
    """

    def cartpole_wrapper(s, t, u):

```

```

        """Helper function to get cartpole() into a form preferred by odeint,
        which expects t as the second arg"""
        return cartpole(s, u)

# PART (c) #####
# INSTRUCTIONS: Complete the function to simulate the cartpole system
# Hint: use the cartpole wrapper above with odeint

# breakpoint()

s = np.zeros((len(t), n))
s[0,:] = s0

u = np.zeros((len(t),m)) # [K@s0] #K @ s
# print("State Variable: ", s_ref[0,:], s0)
# breakpoint()
u[0] = K @ (s0 - s_ref[0,:]).T - u_ref
# breakpoint()
for i, tk in enumerate(t[:-1]):

    # print(i, u[i,0], s[i,:])
    sol = odeint(cartpole_wrapper, s[i,:], t[0:0+2], (u[i],))
    # breakpoint()
    u[i+1] = K @ (sol[1] - s_ref[i]).T - u_ref
    s[i+1,:] = sol[1]

# END PART (c) #####
return s, u

def compute_lti_matrices() -> tuple[np.ndarray, np.ndarray]:
    """Compute the linearized dynamics matrices A and B of the LTI system

    Returns:
        tuple[np.ndarray, np.ndarray]: Tuple of:
            np.ndarray: The A (dynamics) matrix, shape (n, n)
            np.ndarray: The B (controls) matrix, shape (n, m)
    """
# PART (a) #####
df_ds = np.array(
    [
        [0, 0, 1, 0],
        [0, 0, 0, 1],
        [0, mp*g/mc, 0, 0],
        [0, (mc+mp)*g/(mc*L), 0, 0]
    ]
)

```

```

df_du = np.array(
    [
        [0],
        [0],
        [1/mc],
        [1/(mc*L)]
    ]
)

# INSTRUCTIONS: Construct the A and B matrices
A = np.eye(4) + dt * df_ds
B = dt * df_du
# END PART (a) #####
return A, B

def plot_state_and_control_history(
    s: np.ndarray, u: np.ndarray, t: np.ndarray, s_ref: np.ndarray, name: str
) -> None:
    """Helper function for cartpole visualization

    Args:
        s (np.ndarray): State history, shape (num_timesteps, n)
        u (np.ndarray): Control history, shape (num_timesteps, m)
        t (np.ndarray): Times, shape (num_timesteps,)
        s_ref (np.ndarray): Reference state  $\bar{s}$ , evaluated at each time  $t$ .
    Shape (num_timesteps, n)
        name (str): Filename prefix for saving figures
    """
    fig, axes = plt.subplots(1, n + m, dpi=150, figsize=(15, 2))
    plt.subplots_adjust(wspace=0.35)
    labels_s = (r"$x(t)$", r"$\theta(t)$", r"$\dot{x}(t)$",
    ↪r"$\dot{\theta}(t)$")
    labels_u = (r"$u(t)$",)
    for i in range(n):
        axes[i].plot(t, s[:, i])
        axes[i].plot(t, s_ref[:, i], "--")
        axes[i].set_xlabel(r"$t$")
        axes[i].set_ylabel(labels_s[i])
    for i in range(m):
        axes[n + i].plot(t, u[:, i])
        axes[n + i].set_xlabel(r"$t$")
        axes[n + i].set_ylabel(labels_u[i])
    plt.savefig(f"{name}.png", bbox_inches="tight")
    plt.show()

```

```

if animate:
    fig, ani = animate_cartpole(t, s[:, 0], s[:, 1])
    ani.save(f"{name}.mp4", writer="ffmpeg")
    plt.show()

def main():
    # Part A
    A, B = compute_lti_matrices()


    # Part B
    Q = np.eye(n) * 10 # state cost matrix
    R = np.eye(m) # control cost matrix
    K = ricatti_recursion(A, B, Q, R)


    # Part C
    t = np.arange(0.0, 30.0, 1 / 10)
    s_ref = np.array([0.0, np.pi, 0.0, 0.0]) * np.ones((t.size, 1))
    u_ref = np.array([0.0])
    s0 = np.array([0.0, 3*np.pi/4, 0.0, 0.0])
    s, u = simulate(t, s_ref, u_ref, s0, K)
    plot_state_and_control_history(s, u, t, s_ref, "cartpole_balance")

    # Part D
    # Note: t, u_ref unchanged from part c
    s_ref = np.array([reference(ti) for ti in t])
    s0 = np.array([0.0, np.pi, 0.0, 0.0])
    s, u = simulate(t, s_ref, u_ref, s0, K)
    plot_state_and_control_history(s, u, t, s_ref, "cartpole_balance_tv")

if __name__ == "__main__":
    main()

```





0.2 Cart Pole Balancing

```
[1]: """
    Solution code for the problem "Cart-pole balance".

    Autonomous Systems Lab (ASL), Stanford University
    """

import numpy as np
from scipy.integrate import odeint
import matplotlib.pyplot as plt

from animations import animate_cartpole

import pdb

# Constants
n = 4 # state dimension
m = 1 # control dimension
mp = 2.0 # pendulum mass
mc = 10.0 # cart mass
L = 1.0 # pendulum length
g = 9.81 # gravitational acceleration
dt = 0.1 # discretization time step
animate = False # whether or not to animate results
```

```

def cartpole(s: np.ndarray, u: np.ndarray) -> np.ndarray:
    """Compute the cart-pole state derivative

    Args:
        s (np.ndarray): The cartpole state: [x, theta, x_dot, theta_dot], shape (n,)
        u (np.ndarray): The cartpole control: [F_x], shape (m,)

    Returns:
        np.ndarray: The state derivative, shape (n,)
    """
    x, theta, dx, dtheta = s
    sin_theta, cos_theta = np.sin(theta), np.cos(theta)
    h = mc + mp * (sin_theta**2)
    ds = np.array([
        dx,
        dtheta,
        (mp * sin_theta * (L * (dtheta**2) + g * cos_theta) + u[0]) / h,
        -((mc + mp) * g * sin_theta + mp * L * (dtheta**2) * sin_theta * cos_theta + u[0] * cos_theta) / (h * L),
    ])
    return ds

def reference(t: float) -> np.ndarray:
    """Compute the reference state (s_bar) at time t

    Args:
        t (float): Evaluation time

    Returns:
        np.ndarray: Reference state, shape (n,)
    """
    a = 10.0 # Amplitude
    T = 10.0 # Period
    # breakpoint()
    # PART (d) #####
    # INSTRUCTIONS: Compute the reference state for a given time
    # raise NotImplementedError()
    return np.array([a * np.sin(2*np.pi*t/T), np.pi, 2*np.pi*a/T * np.cos(2*np.pi*t/T), 0]).T
    # END PART (d) #####

```



```

def ricatti_recursion(
    A: np.ndarray, B: np.ndarray, Q: np.ndarray, R: np.ndarray
) -> np.ndarray:
    """Compute the gain matrix K through Ricatti recursion

    Args:
        A (np.ndarray): Dynamics matrix, shape (n, n)
        B (np.ndarray): Controls matrix, shape (n, m)
        Q (np.ndarray): State cost matrix, shape (n, n)
        R (np.ndarray): Control cost matrix, shape (m, m)

    Returns:
        np.ndarray: Gain matrix K, shape (m, n)
    """
    eps = 1e-4 # Riccati recursion convergence tolerance
    max_iters = 1000 # Riccati recursion maximum number of iterations
    P_prev = np.zeros((n, n)) # initialization
    converged = False
    for i in range(max_iters):
        # PART (b) #####
        # INSTRUCTIONS: Apply the Ricatti equation until convergence
        K = -np.linalg.inv(R + B.T @ P_prev @ B) @ B.T @ P_prev @ A
        P_k = Q + A.T @ P_prev @ (A + B @ K)

        # termination condition
        if np.max(np.abs(P_prev - P_k)) < 1e-4:
            converged = True
            break
        else:
            P_prev = P_k
        # END PART (b) #####
    if not converged:
        raise RuntimeError("Ricatti recursion did not converge!")
    print("K:", K)
    return K


def simulate(
    t: np.ndarray, s_ref: np.ndarray, u_ref: np.ndarray, s0: np.ndarray, K: np.
    ↪ ndarray
) -> tuple[np.ndarray, np.ndarray]:
    """Simulate the cartpole

    Args:
        t (np.ndarray): Evaluation times, shape (num_timesteps,)
        s_ref (np.ndarray): Reference state s_bar, evaluated at each time t.
        ↪ Shape (num_timesteps, n)

```

```

    u_ref (np.ndarray): Reference control  $\bar{u}$ , shape (m,)
    s0 (np.ndarray): Initial state, shape (n,)
    K (np.ndarray): Feedback gain matrix (Ricatti recursion result), shape  $\hookrightarrow$ 
     $\hookrightarrow$  (m, n)

```

Returns:

```

    tuple[np.ndarray, np.ndarray]: Tuple of:
        np.ndarray: The state history, shape (num_timesteps, n)
        np.ndarray: The control history, shape (num_timesteps, m)
    """

```

```

def cartpole_wrapper(s, t, u):
    """Helper function to get cartpole() into a form preferred by odeint,  $\hookrightarrow$ 
     $\hookrightarrow$  which expects t as the second arg"""
    return cartpole(s, u)

```

```

# PART (c) #####
# INSTRUCTIONS: Complete the function to simulate the cartpole system
# Hint: use the cartpole wrapper above with odeint

```

```

# breakpoint()

```

```

s = np.zeros((len(t), n))
s[0,:] = s0

```

```

u = np.zeros((len(t),m)) # [K@s0] #K @ s
# print("State Variable: ", s_ref[0,:], s0)
# breakpoint()
u[0] = K @ (s0 - s_ref[0,:]).T - u_ref
# breakpoint()
for i, tk in enumerate(t[:-1]):

```

```

    # print(i, u[i,0], s[i,:])
    sol = odeint(cartpole_wrapper, s[i,:], t[0:0+2], (u[i],))
    # breakpoint()
    u[i+1] = K @ (sol[1] - s_ref[i]).T - u_ref
    s[i+1,:] = sol[1]

```

```

# END PART (c) #####
return s, u

```

```

def compute_lti_matrices() -> tuple[np.ndarray, np.ndarray]:
    """Compute the linearized dynamics matrices A and B of the LTI system

```

Returns:

```

    tuple[np.ndarray, np.ndarray]: Tuple of:

```

```

        np.ndarray: The A (dynamics) matrix, shape (n, n)
        np.ndarray: The B (controls) matrix, shape (n, m)
    """
    # PART (a) #####
    df_ds = np.array(
        [
            [0, 0, 1, 0],
            [0, 0, 0, 1],
            [0, mp*g/mc, 0, 0],
            [0, (mc+mp)*g/(mc*L), 0, 0]
        ]
    )

    df_du = np.array(
        [
            [0],
            [0],
            [1/mc],
            [1/(mc*L)]
        ]
    )

    # INSTRUCTIONS: Construct the A and B matrices
    A = np.eye(4) + dt * df_ds
    B = dt * df_du
    # END PART (a) #####
    return A, B

def plot_state_and_control_history(
    s: np.ndarray, u: np.ndarray, t: np.ndarray, s_ref: np.ndarray, name: str
) -> None:
    """Helper function for cartpole visualization

    Args:
        s (np.ndarray): State history, shape (num_timesteps, n)
        u (np.ndarray): Control history, shape (num_timesteps, m)
        t (np.ndarray): Times, shape (num_timesteps,)
        s_ref (np.ndarray): Reference state s_bar, evaluated at each time t.
    Shape (num_timesteps, n)
    name (str): Filename prefix for saving figures
    """
    fig, axes = plt.subplots(1, n + m, dpi=150, figsize=(15, 2))
    plt.subplots_adjust(wspace=0.35)
    labels_s = (r"$x(t)$", r"$\theta(t)$", r"$\dot{x}(t)$",
    ↪r"$\dot{\theta}(t)$")
    labels_u = (r"$u(t)$",)

```

```

for i in range(n):
    axes[i].plot(t, s[:, i])
    axes[i].plot(t, s_ref[:, i], "--")
    axes[i].set_xlabel(r"$t$")
    axes[i].set_ylabel(labels_s[i])
for i in range(m):
    axes[n + i].plot(t, u[:, i])
    axes[n + i].set_xlabel(r"$t$")
    axes[n + i].set_ylabel(labels_u[i])
plt.savefig(f"{name}.png", bbox_inches="tight")
plt.show()

if animate:
    fig, ani = animate_cartpole(t, s[:, 0], s[:, 1])
    ani.save(f"{name}.mp4", writer="ffmpeg")
    plt.show()

def main():
    # Part A
    A, B = compute_lti_matrices()

    # Part B
    Q = np.eye(n) * 10 # state cost matrix
    R = np.eye(m) # control cost matrix
    K = ricatti_recursion(A, B, Q, R)

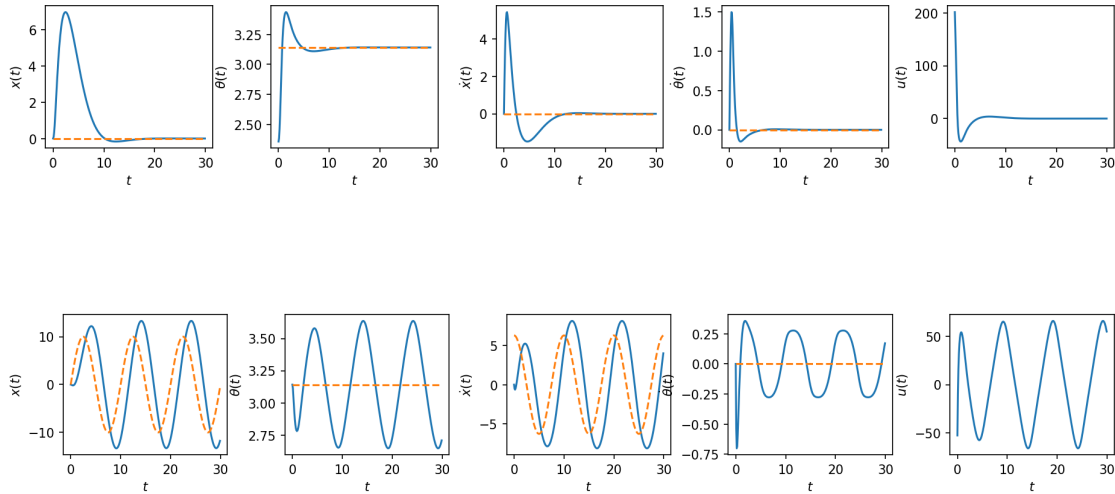
    # Part C
    t = np.arange(0.0, 30.0, 1 / 10)
    s_ref = np.array([0.0, np.pi, 0.0, 0.0]) * np.ones((t.size, 1))
    u_ref = np.array([0.0])
    s0 = np.array([0.0, 3*np.pi/4, 0.0, 0.0])
    s, u = simulate(t, s_ref, u_ref, s0, K)
    plot_state_and_control_history(s, u, t, s_ref, "cartpole_balance")

    # Part D
    # Note: t, u_ref unchanged from part c
    s_ref = np.array([reference(ti) for ti in t])
    s0 = np.array([0.0, np.pi, 0.0, 0.0])
    s, u = simulate(t, s_ref, u_ref, s0, K)
    plot_state_and_control_history(s, u, t, s_ref, "cartpole_balance_tv")

if __name__ == "__main__":
    main()

```

K: [[2.26379564 -256.34294732 8.42679164 -76.04341843]]



```
[2]: """
Solution code for the problem "Cart-pole balance".

Autonomous Systems Lab (ASL), Stanford University
"""

import numpy as np
from scipy.integrate import odeint
import matplotlib.pyplot as plt

from animations import animate_cartpole

import pdb

# Constants
n = 4 # state dimension
m = 1 # control dimension
mp = 2.0 # pendulum mass
mc = 10.0 # cart mass
L = 1.0 # pendulum length
g = 9.81 # gravitational acceleration
dt = 0.1 # discretization time step
animate = False # whether or not to animate results

def cartpole(s: np.ndarray, u: np.ndarray) -> np.ndarray:
    """Compute the cart-pole state derivative

    Args:
```

```

    s (np.ndarray): The cartpole state: [x, theta, x_dot, theta_dot], shape (n,)
    u (np.ndarray): The cartpole control: [F_x], shape (m,)

Returns:
    np.ndarray: The state derivative, shape (n,)
    """
    x, theta, dx, d = s
    sin_theta, cos_theta = np.sin(theta), np.cos(theta)
    h = mc + mp * (sin_theta**2)
    ds = np.array(
        [
            dx,
            d,
            (mp * sin_theta * (L * (d**2) + g * cos_theta) + u[0]) / h,
            -((mc + mp) * g * sin_theta + mp * L * (d**2) * sin_theta * cos_theta + u[0] * cos_theta) / (h * L),
        ]
    )
    return ds

def reference(t: float) -> np.ndarray:
    """Compute the reference state (s_bar) at time t

    Args:
        t (float): Evaluation time

    Returns:
        np.ndarray: Reference state, shape (n,)
        """
    a = 10.0 # Amplitude
    T = 10.0 # Period
    # breakpoint()
    # PART (d) #####
    # INSTRUCTIONS: Compute the reference state for a given time
    # raise NotImplementedError()
    return np.array([a * np.sin(2*np.pi*t/T), np.pi, 2*np.pi*a/T * np.cos(2*np.pi*t/T), 0]).T
    # END PART (d) #####

def ricatti_recursion(
    A: np.ndarray, B: np.ndarray, Q: np.ndarray, R: np.ndarray
) -> np.ndarray:
    """Compute the gain matrix K through Ricatti recursion

```

Args:

A (np.ndarray): Dynamics matrix, shape (n, n)
B (np.ndarray): Controls matrix, shape (n, m)
Q (np.ndarray): State cost matrix, shape (n, n)
R (np.ndarray): Control cost matrix, shape (m, m)

Returns:

np.ndarray: Gain matrix K, shape (m, n)

```
"""
eps = 1e-4 # Riccati recursion convergence tolerance
max_iters = 1000 # Riccati recursion maximum number of iterations
P_prev = np.zeros((n, n)) # initialization
converged = False
for i in range(max_iters):
    # PART (b) #####
    # INSTRUCTIONS: Apply the Ricatti equation until convergence
    K = -np.linalg.inv(R + B.T @ P_prev @ B) @ B.T @ P_prev @ A
    P_k = Q + A.T @ P_prev @ (A + B @ K)

    # termination condition
    if np.max(np.abs(P_prev-P_k)) < 1e-4:
        converged = True
        break
    else:
        P_prev = P_k
    # END PART (b) #####
if not converged:
    raise RuntimeError("Ricatti recursion did not converge!")
print("K:", K)
return K
```

```
def simulate(
    t: np.ndarray, s_ref: np.ndarray, u_ref: np.ndarray, s0: np.ndarray, K: np.
    ↪ ndarray
) -> tuple[np.ndarray, np.ndarray]:
    """Simulate the cartpole
```

Args:

t (np.ndarray): Evaluation times, shape (num_timesteps,)
s_ref (np.ndarray): Reference state s_bar, evaluated at each time t. ↵
↪ Shape (num_timesteps, n)
u_ref (np.ndarray): Reference control u_bar, shape (m,)
s0 (np.ndarray): Initial state, shape (n,)
K (np.ndarray): Feedback gain matrix (Ricatti recursion result), shape ↵
↪ (m, n)

```

Returns:
    tuple[np.ndarray, np.ndarray]: Tuple of:
        np.ndarray: The state history, shape (num_timesteps, n)
        np.ndarray: The control history, shape (num_timesteps, m)
    """

def cartpole_wrapper(s, t, u):
    """Helper function to get cartpole() into a form preferred by odeint,
    which expects t as the second arg"""
    return cartpole(s, u)

# PART (c) #####
# INSTRUCTIONS: Complete the function to simulate the cartpole system
# Hint: use the cartpole wrapper above with odeint

# breakpoint()

s = np.zeros((len(t), n))
s[0,:] = s0

u = np.zeros((len(t),m)) # [K@s0] #K @ s
# print("State Variable: ", s_ref[0,:], s0)
# breakpoint()
u[0] = K @ (s0 - s_ref[0,:]).T - u_ref
# breakpoint()
for i, tk in enumerate(t[:-1]):

    # print(i, u[i,0], s[i,:])
    sol = odeint(cartpole_wrapper, s[i,:], t[0:0+2], (u[i],))
    # breakpoint()
    u[i+1] = K @ (sol[1] - s_ref[i]).T - u_ref
    s[i+1,:] = sol[1]

# END PART (c) #####
return s, u

def compute_lti_matrices() -> tuple[np.ndarray, np.ndarray]:
    """Compute the linearized dynamics matrices A and B of the LTI system

    Returns:
        tuple[np.ndarray, np.ndarray]: Tuple of:
            np.ndarray: The A (dynamics) matrix, shape (n, n)
            np.ndarray: The B (controls) matrix, shape (n, m)
        """

    # PART (a) #####
    df_ds = np.array(

```



```

        [
            [0, 0, 1, 0],
            [0, 0, 0, 1],
            [0, mp*g/mc, 0, 0],
            [0, (mc+mp)*g/(mc*L), 0, 0]
        ]
    )

    df_du = np.array(
        [
            [0],
            [0],
            [1/mc],
            [1/(mc*L)]
        ]
    )

    # INSTRUCTIONS: Construct the A and B matrices
    A = np.eye(4) + dt * df_ds
    B = dt * df_du
    # END PART (a) #####
    return A, B

def plot_state_and_control_history(
    s: np.ndarray, u: np.ndarray, t: np.ndarray, s_ref: np.ndarray, name: str
) -> None:
    """Helper function for cartpole visualization

    Args:
        s (np.ndarray): State history, shape (num_timesteps, n)
        u (np.ndarray): Control history, shape (num_timesteps, m)
        t (np.ndarray): Times, shape (num_timesteps,)
        s_ref (np.ndarray): Reference state s_bar, evaluated at each time t.
    Shape (num_timesteps, n)
        name (str): Filename prefix for saving figures
    """
    fig, axes = plt.subplots(1, n + m, dpi=150, figsize=(15, 2))
    plt.subplots_adjust(wspace=0.35)
    labels_s = (r"$x(t)$", r"$\theta(t)$", r"$\dot{x}(t)$",
    ↪r"$\dot{\theta}(t)$")
    labels_u = (r"$u(t)$",)
    for i in range(n):
        axes[i].plot(t, s[:, i])
        axes[i].plot(t, s_ref[:, i], "--")
        axes[i].set_xlabel(r"$t$")
        axes[i].set_ylabel(labels_s[i])

```

```

for i in range(m):
    axes[n + i].plot(t, u[:, i])
    axes[n + i].set_xlabel(r"$t$")
    axes[n + i].set_ylabel(labels_u[i])
plt.savefig(f"{name}.png", bbox_inches="tight")
plt.show()

if animate:
    fig, ani = animate_cartpole(t, s[:, 0], s[:, 1])
    ani.save(f"{name}.mp4", writer="ffmpeg")
    plt.show()

def main():
    # Part A
    A, B = compute_lti_matrices()

    # Part B
    Q = np.eye(n) * 1 # state cost matrix
    R = np.eye(m) # control cost matrix
    K = ricatti_recursion(A, B, Q, R)

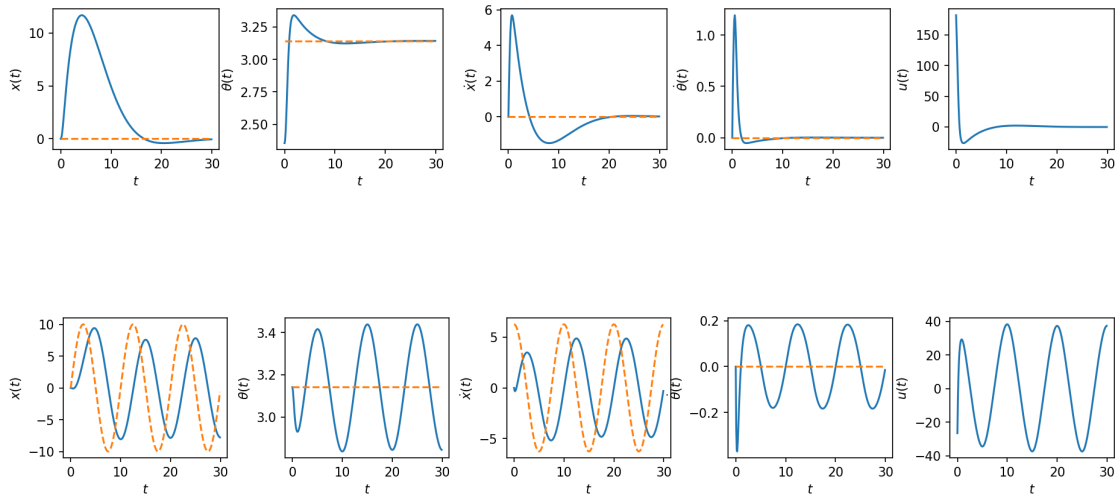
    # Part C
    t = np.arange(0.0, 30.0, 1 / 10)
    s_ref = np.array([0.0, np.pi, 0.0, 0.0]) * np.ones((t.size, 1))
    u_ref = np.array([0.0])
    s0 = np.array([0.0, 3*np.pi/4, 0.0, 0.0])
    s, u = simulate(t, s_ref, u_ref, s0, K)
    plot_state_and_control_history(s, u, t, s_ref, "cartpole_balance")

    # Part D
    # Note: t, u_ref unchanged from part c
    s_ref = np.array([reference(ti) for ti in t])
    s0 = np.array([0.0, np.pi, 0.0, 0.0])
    s, u = simulate(t, s_ref, u_ref, s0, K)
    plot_state_and_control_history(s, u, t, s_ref, "cartpole_balance_tv")

if __name__ == "__main__":
    main()

```

K: [[0.7291397 -231.85419273 4.21967187 -68.24742822]]



1 Cart Pole Balance

```
[ ]: """
    Starter code for the problem "Cart-pole swing-up".

    Autonomous Systems Lab (ASL), Stanford University
    """

import time

from animations import animate_cartpole

import jax
import jax.numpy as jnp

import matplotlib.pyplot as plt

import numpy as np

from scipy.integrate import odeint

def linearize(f, s, u):
    """Linearize the function `f(s, u)` around `(s, u)`.

    Arguments
    -----
    f : callable
        A nonlinear function with call signature `f(s, u)`.
```

```

s : numpy.ndarray
    The state (1-D).
u : numpy.ndarray
    The control input (1-D).

Returns
-----
A : numpy.ndarray
    The Jacobian of `f` at `(s, u)`, with respect to `s`.
B : numpy.ndarray
    The Jacobian of `f` at `(s, u)`, with respect to `u`.
"""
# WRITE YOUR CODE BELOW #####
# INSTRUCTIONS: Use JAX to compute `A` and `B` in one line.
# raise NotImplementedError()
A, B = jax.jacrev(f, argnums=(0, 1))(s, u)
#####
return A, B

def ilqr(f, s0, s_goal, N, Q, R, QN, eps=1e-3, max_iters=1000):
    """Compute the iLQR set-point tracking solution.

    Arguments
    -----
    f : callable
        A function describing the discrete-time dynamics, such that
         $s[k+1] = f(s[k], u[k])$ .
    s0 : numpy.ndarray
        The initial state (1-D).
    s_goal : numpy.ndarray
        The goal state (1-D).
    N : int
        The time horizon of the LQR cost function.
    Q : numpy.ndarray
        The state cost matrix (2-D).
    R : numpy.ndarray
        The control cost matrix (2-D).
    QN : numpy.ndarray
        The terminal state cost matrix (2-D).
    eps : float, optional
        Termination threshold for iLQR.
    max_iters : int, optional
        Maximum number of iLQR iterations.

    Returns
    -----

```

```

s_bar : numpy.ndarray
    A 2-D array where `s_bar[k]` is the nominal state at time step `k`,
    for `k = 0, 1, ..., N-1`
u_bar : numpy.ndarray
    A 2-D array where `u_bar[k]` is the nominal control at time step `k`,
    for `k = 0, 1, ..., N-1`
Y : numpy.ndarray
    A 3-D array where `Y[k]` is the matrix gain term of the iLQR control
    law at time step `k`, for `k = 0, 1, ..., N-1`
y : numpy.ndarray
    A 2-D array where `y[k]` is the offset term of the iLQR control law
    at time step `k`, for `k = 0, 1, ..., N-1`
"""
if max_iters <= 1:
    raise ValueError("Argument `max_iters` must be at least 1.")
n = Q.shape[0] # state dimension
m = R.shape[0] # control dimension

# Initialize gains `Y` and offsets `y` for the policy
Y = np.zeros((N, m, n))
y = np.zeros((N, m))

def rollout(x0, u_trj):
    x_trj = np.zeros((u_trj.shape[0] + 1, x0.shape[0]))
    # TODO: Define the rollout here and return the state trajectory x_trj:
    ↪ [N, number of states]
    x_trj[0] = x0
    for i, u in enumerate(u_trj):
        x_trj[i+1] = fd(x_trj[i], u)
    return x_trj

# Initialize the nominal trajectory `(s_bar, u_bar)`, and the
# deviations `(ds, du)`
u_bar = np.random.rand(N, m)
# s_bar[0] = s0
s_bar = rollout(s0, u_bar)
for k in range(N):
    s_bar[k + 1] = f(s_bar[k], u_bar[k])
ds = np.zeros((N + 1, n))
du = np.zeros((N, m))

# Define Helper Function
def forward_pass(x_trj, u_trj, k_trj, K_trj):
    x_trj_new = np.zeros(x_trj.shape)
    x_trj_new[0, :] = x_trj[0, :]
    u_trj_new = np.zeros(u_trj.shape)
    # TODO: Implement the forward pass here

```

```

        for n in range(u_trj.shape[0]):
            # Note, converting from deviation variable to actual value
            ↪variable
            u_trj_new[n,:] = u_trj[n,:] + k_trj[n,:] + K_trj[n,:] @
            ↪(x_trj_new[n,:]-x_trj[n,:])# Apply feedback law
            x_trj_new[n+1,:] = fd(x_trj_new[n,:], u_trj_new[n,:]) # Apply
            ↪dynamics
            return x_trj_new, u_trj_new

    @jax.jit
    def cost_stage(x, u):
        return 1/2 * ((x - s_goal).T @ Q @ (x-s_goal) + u.T @ R @ u)

    @jax.jit
    def stage(x, u):
        l = cost_stage
        l_x = jax.jacrev(l, argnums=0)
        l_u = jax.jacrev(l, argnums=1)
        l_xx = jax.jacrev(l_x, argnums=0)
        l_ux = jax.jacrev(l_u, argnums=0)
        l_uu = jax.jacrev(l_u, argnums=1)

        # f = fd
        f_x = jax.jacrev(fd, argnums=0)
        f_u = jax.jacrev(fd, argnums=1)

        return l_x(x, u), l_u(x, u), l_xx(x, u), l_ux(x, u), l_uu(x, u), f_x(x,
            ↪u), f_u(x, u)

    @jax.jit
    def cost_final(x):
        return 1/2 * (x - s_goal).T @ QN @ (x-s_goal)

    @jax.jit
    def final(x):
        l_final = cost_final
        l_final_x = jax.jacrev(l_final, argnums=0)
        l_final_xx = jax.jacrev(l_final_x, argnums=0)

        return l_final_x(x), l_final_xx(x)

    @jax.jit
    def cost_trj(x_trj, u_trj):
        total = 0.0
        total = (
            cost_final(x_trj[-1]))+

```

```

        jnp.sum(jnp.array([cost_stage(x, u) for x, u in zip(x_trj[:-1], u
↪u_trj)]))
    )

    return total

@jax.jit
def gains(Q_uu, Q_u, Q_ux):
    Q_uu_inv = jnp.linalg.inv(Q_uu)
    # TODO: Implement the feedforward gain k and feedback gain K.
    k = - Q_uu_inv @ Q_u.T #np.zeros(Q_u.shape)
    K = - Q_uu_inv @ Q_ux #np.zeros(Q_ux.shape)
    return k, K

@jax.jit
def V_terms(Q_x, Q_u, Q_xx, Q_ux, Q_uu, K, k):
    # TODO: Implement V_x and V_xx, hint: use the A.dot(B) function for
↪matrix multiplication.
    V_x = Q_x + K.T @ Q_u + k.T @ Q_ux + K.T @ Q_uu @ k #np.zeros(Q_x.shape)
    # print(Q_xx.shape, Q_ux.T.shape, K.shape, K.T.shape, Q_ux.shape, K.
↪shape) #np.zeros(Q_xx.shape)
    V_xx = Q_xx + 2 * Q_ux.T @ K + K.T @ Q_uu @ K #np.zeros(Q_xx.shape)

    return V_x, V_xx

@jax.jit
def Q_terms(l_x, l_u, l_xx, l_ux, l_uu, f_x, f_u, V_x, V_xx):
    # TODO: Define the Q-terms here
    Q_x = l_x + V_x.T @ f_x #np.zeros(l_x.shape)
    Q_u = l_u + V_x.T @ f_u #np.zeros(l_u.shape)
    Q_xx = l_xx + f_x.T @ V_xx @ f_x #np.zeros(l_xx.shape)
    Q_ux = l_ux + f_u.T @ V_xx @ f_x #np.zeros(l_ux.shape)
    Q_uu = l_uu + f_u.T @ V_xx @ f_u #np.zeros(l_uu.shape)
    return Q_x, Q_u, Q_xx, Q_ux, Q_uu

def backward_pass(x_trj, u_trj, regu=0):
    k_trj = np.zeros([u_trj.shape[0], u_trj.shape[1]])
    K_trj = np.zeros([u_trj.shape[0], u_trj.shape[1], x_trj.shape[1]])
    # expected_cost_redu = 0
    # TODO: Set terminal boundary condition here (V_x, V_xx)
    V_x, V_xx = final(x_trj[-1])

    for n in range(u_trj.shape[0] - 1, -1, -1):
        # TODO: First compute derivatives, then the Q-terms
        l_x, l_u, l_xx, l_ux, l_uu, f_x, f_u = stage(x_trj[n], u_trj[n])

```

```

        Q_x, Q_u, Q_xx, Q_ux, Q_uu = Q_terms(l_x, l_u, l_xx, l_ux, l_uu,
↪f_x, f_u, V_x, V_xx)

        # We add regularization to ensure that Q_uu is invertible and
↪nicely conditioned
        Q_uu_regu = Q_uu + np.eye(Q_uu.shape[0]) * regu
        k, K = gains(Q_uu_regu, Q_u, Q_ux)
        k_trj[n, :] = k
        K_trj[n, :, :] = K
        V_x, V_xx = V_terms(Q_x, Q_u, Q_xx, Q_ux, Q_uu, K, k)
        # expected_cost_redu += expected_cost_reduction(Q_u, Q_uu, k)
    return k_trj, K_trj#, expected_cost_redu

# iLQR loop
converged = False
cost = np.inf
for _ in range(max_iters):
    # Linearize the dynamics at each step `k` of `(s_bar, u_bar)`
    A, B = jax.vmap(linearize, in_axes=(None, 0, 0))(f, s_bar[:-1], u_bar)
    A, B = np.array(A), np.array(B)

    # PART (c) #####
    # INSTRUCTIONS: Update `Y`, `y`, `ds`, `du`, `s_bar`, and `u_bar`.
    # raise NotImplementedError()

    y, Y = backward_pass(s_bar, u_bar, regu=0)

    s_bar_new, u_bar_new = forward_pass(s_bar, u_bar, y, Y)

    print(cost_trj(s_bar_new, u_bar_new))

    du = u_bar_new - u_bar
    s_bar = s_bar_new
    u_bar = u_bar_new

    #####

    if np.max(np.abs(du)) < eps:
        converged = True
        print('Converged')
        break
    # print('one more')
if not converged:
    print("iLQR did not converge!")
return s_bar, u_bar, Y, y

```



```

def cartpole(s, u):
    """Compute the cart-pole state derivative."""
    mp = 2.0 # pendulum mass
    mc = 10.0 # cart mass
    L = 1.0 # pendulum length
    g = 9.81 # gravitational acceleration

    x, , dx, d = s
    sin, cos = jnp.sin(), jnp.cos()
    h = mc + mp * (sin**2)
    ds = jnp.array(
        [
            dx,
            d,
            (mp * sin * (L * (d**2) + g * cos) + u[0]) / h,
            -((mc + mp) * g * sin + mp * L * (d**2) * sin * cos + u[0] * cos)
              / (h * L),
        ]
    )
    return ds

# Define constants
n = 4 # state dimension
m = 1 # control dimension
Q = np.diag(np.array([10.0, 10.0, 2.0, 2.0])) # state cost matrix
R = 1e-2 * np.eye(m) # control cost matrix
QN = 1e2 * np.eye(n) # terminal state cost matrix
s0 = np.array([0.0, 0.0, 0.0, 0.0]) # initial state
s_goal = np.array([0.0, np.pi, 0.0, 0.0]) # goal state
T = 10.0 # simulation time
dt = 0.1 # sampling time
animate = False # flag for animation
closed_loop = False # flag for closed-loop control

# Initialize continuous-time and discretized dynamics
f = jax.jit(cartpole)
fd = jax.jit(lambda s, u, dt=dt: s + dt * f(s, u))

# Compute the iLQR solution with the discretized dynamics
print("Computing iLQR solution ... ", end="", flush=True)
start = time.time()
t = np.arange(0.0, T, dt)
N = t.size - 1
s_bar, u_bar, Y, y = ilqr(fd, s0, s_goal, N, Q, R, QN)
print("done! ({:.2f} s)".format(time.time() - start), flush=True)

```

```

# Plot iLQR solution
# fig, axes = plt.subplots(1, n + m, dpi=150, figsize=(15, 2))
# plt.subplots_adjust(wspace=0.45)
# labels_s = (r"$x(t)$", r"$\theta(t)$", r"$\dot{x}(t)$", r"$\dot{\theta}(t)$")
# labels_u = (r"$u(t)$",)
# for i in range(n):
#     axes[i].plot(t, s_bar[:, i])
#     axes[i].set_xlabel(r"$t$")
#     axes[i].set_ylabel(labels_s[i])
# for i in range(m):
#     axes[n + i].plot(t[:-1], u_bar[:, i])
#     axes[n + i].set_xlabel(r"$t$")
#     axes[n + i].set_ylabel(labels_u[i])

# axes[1].axhline(np.pi, linestyle="--", color="tab:orange")

# Simulate on the true continuous-time system
print("Simulating ... ", end="", flush=True)
start = time.time()
s = np.zeros((N + 1, n))
u = np.zeros((N, m))
s[0] = s0
for k in range(N):
    # PART (d) #####
    # INSTRUCTIONS: Compute either the closed-loop or open-loop value of
    # `u[k]`, depending on the Boolean flag `closed_loop`.
    if closed_loop:
        u[k] = 0.0
        raise NotImplementedError()
    else: # do open-loop control
        u[k] = u_bar[k]
        # raise NotImplementedError()
    #####
    s[k + 1] = odeint(lambda s, t: f(s, u[k]), s[k], t[k : k + 2])[1]
print("done! ({:.2f} s)".format(time.time() - start), flush=True)

# Plot
fig, axes = plt.subplots(1, n + m, dpi=150, figsize=(15, 2))
plt.subplots_adjust(wspace=0.45)
labels_s = (r"$x(t)$", r"$\theta(t)$", r"$\dot{x}(t)$", r"$\dot{\theta}(t)$")
labels_u = (r"$u(t)$",)
for i in range(n):
    axes[i].plot(t, s[:, i])
    axes[i].set_xlabel(r"$t$")
    axes[i].set_ylabel(labels_s[i])
for i in range(m):

```

```

    axes[n + i].plot(t[:-1], u[:, i])
    axes[n + i].set_xlabel(r"$t$")
    axes[n + i].set_ylabel(labels_u[i])
if closed_loop:
    plt.savefig("cartpole_swingup_cl.png", bbox_inches="tight")
else:
    plt.savefig("cartpole_swingup_ol.png", bbox_inches="tight")
plt.show()

if animate:
    fig, ani = animate_cartpole(t, s[:, 0], s[:, 1])
    ani.save("cartpole_swingup.mp4", writer="ffmpeg")
    plt.show()

```

[]: